

Nils Lukas

Assistant Professor • MBZUAI • Abu Dhabi, UAE

✉ nils.lukas@mbzuai.ac.ae • [nillslukas.github.io](https://github.com/nillslukas) • [Google Scholar](#)

Updated on July 26, 2026

Research Interests

Design secure and private AI agents in the presence of untrustworthy **providers**, **data**, **models**, and **users**.

1. **Providers**: Confidential computing & encrypted inference.
2. **Data**: Poisoning, prompt injection & training-data privacy.
3. **Models**: Scalable oversight & differentially private inference.
4. **Users**: Detect misuse via watermarking & safety monitoring.

Research Grants

Awarded	\$1,458,000
[United AI-Saqer Group] MedBoard: Privacy-preserving Multi-Agent AI for Healthcare	\$136,000 · 2026
PI: Nils Lukas · Co-PIs: Praneeth Vepakomma , Preslav Nakov	
[IAAI Research Program] Trust Layer for Safe Smallholder Advisory	\$500,000 · 2026
PI: Salem Lahlou · Co-PIs: Nils Lukas, Martin Takáč	
[Amazon Research Awards – Special Call] AdvSim2Real: Robustifying Agents in Simulated Adversarial Environments	\$100,000 · 2025
PI: Nils Lukas · Co-PI: Praneeth Vepakomma	
[Etihad Airways] Conversational Booking Agents	\$450,000 · 2025
PI: Nils Lukas · Co-PIs: Salem Lahlou , Alham Fikri Aji , Martin Takáč , Mingming Gong	
[United AI-Saqer Group] Privacy-preserving Brain–Computer Interfaces	\$136,000 · 2025
PI: Abdulrahman Mahmoud · Co-PIs: Nils Lukas, Elizabeth Churchill	
[TII Funding] GFlowNets for Fuzzing of Agentic Applications	\$136,000 · 2025
PIs: Salem Lahlou & Nils Lukas	

Conference Publications

Underlined authors are students and postdocs I supervise as primary advisor.

25	CopyShield: A Cross-Level Benchmark of Copyright Defenses in LLMs
[EMNLP'26]	Maryam Alshehryari , Dushyant Chauhan , Samuele Poppi , Martin Takáč, Salem Lahlou, Nils Lukas. Conference on Empirical Methods in Natural Language Processing (Main Conference), 2026.
Main Conference AR: 15.4% (2,719/17,669)	
24	TRACER: Trace-Regularized Training Against Indirect Prompt Injection
[EMNLP'26]	Kshitij Mishra , Salem Lahlou, Nils Lukas. Findings of the Association for Computational Linguistics: EMNLP, 2026.
Findings AR: 14.3% (2,533/17,669)	
23	SPQR: A Multi-Dimensional Benchmark for Safety Alignment under Benign Model Adaptation
[ECCV'26]	Mohammed Talha Alam, Nada Saadi, Fahad Shamshad, Nils Lukas, Karthik Nandakumar, Fakhri Karray, Samuele Poppi . European Conference on Computer Vision, 2026.
22	RAVEN: Erasing Invisible Watermarks via Novel View Synthesis
[CVPR'26]	Fahad Shamshad, Nils Lukas, Karthik Nandakumar. IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2026.
AR: 25.4% (4,090/16,092)	
🏆 Oral (top 0.8%)	

21	Collaborative Threshold Watermarking <u>Tameem Bakr</u> , Anish Ambreth, Nils Lukas . The Forty-Third International Conference on Machine Learning, 2026.
[ICML'26] AR: 26.6% (6,352/23,918)	
20	CoRe: Collaborative Reasoning via Cross Teaching <u>Kshitij Mishra</u> , Mirat Aubakirov, Martin Takáč, Nils Lukas , Salem Lahlou. The Forty-Third International Conference on Machine Learning, 2026.
[ICML'26] AR: 26.6% (6,352/23,918)	
19	Forest Before Trees: Latent Superposition for Efficient Visual Reasoning Yubo Wang, Juntian Zhang, Yichen Wu, Yankai Lin, Nils Lukas , <u>Yuhan Liu</u> . The 64th Annual Meeting of the Association for Computational Linguistics, 2026.
[ACL'26] Main Conference AR: 18.9% (2,296/12,148)	
18	Detecting AI-Generated Video: A Vision-Language Dual-View Survey Dylan Xinming Hou, Juntian Zhang, Xu Gu, Yichen Wu, Nils Lukas , Gus Xia, Xiuying Chen, <u>Yuhan Liu</u> . Findings of the Association for Computational Linguistics, 2026.
[ACL'26] Findings AR: 17.8% (2,163/12,148)	
17	DP-Fusion: Token-Level Differentially Private Inference for Large Language Models <u>Rushil Thareja</u> , Preslav Nakov, Praneeth Vepakomma , Nils Lukas . The Fourteenth International Conference on Learning Representations, 2026.
[ICLR'26] AR: 27.4% (5,355/19,525)	
16	SD-E²: Semantic Exploration for Reasoning Under Token Budgets <u>Kshitij Mishra</u> , Nils Lukas , Salem Lahlou. Findings of the European Chapter of the Association for Computational Linguistics, 2026.
[EACL'26] Findings	
15	Mask Image Watermarking Runyi Hu, Jie Zhang, Shiqian Zhao, Nils Lukas , Jiwei Li, Qing Guo, Han Qiu, Tianwei Zhang. The Thirty-Ninth Annual Conference on Neural Information Processing Systems, 2025.
[NeurIPS'25] AR: 24.5% (5,290/21,575)	
14	SPIRIT: Patching Speech Language Models against Jailbreak Attacks Amir Djanibekov, Nurdaulet Mukhituly, Kentaro Inui, Hanan Aldarmaki, Nils Lukas . Empirical Methods in Natural Language Processing (Main Conference), 2025.
[EMNLP'25] AR: 22.2% (1,811/8,172)	
13	Optimizing Adaptive Attacks against Content Watermarks for Language Models Abdulrahman Diaa, <u>Toluwani Aremu</u> , Nils Lukas . The Forty-Second International Conference on Machine Learning, 2025.
[ICML'25] AR: 26.9% (3,260/12,107) ★ Spotlight (Top 2.6%)	
12	Cowpox: Towards the Immunity of VLM-based Multi-Agent Systems Yutong Wu, Jie Zhang, Yiming Li, Chao Zhang, Qing Guo, Han Qiu, Nils Lukas , Tianwei Zhang. The Forty-Second International Conference on Machine Learning, 2025.
[ICML'25] AR: 26.9% (3,260/12,107)	
11	PEPSI: Practically Efficient Private Set Intersection in the Unbalanced Setting Rasoul Mahdavi, Nils Lukas , Faezeh Ebrahimianghazani, Thomas Humphries, Bailey Kacsmar, John Premkumar, Xinda Li, Simon Oya, Ehsan Amjadian, Florian Kerschbaum. In the 33rd USENIX Security Symposium, 2024.
[USENIX'24] AR: 19.1% (417/2,176)	

10	Fast and Private Inference of Deep Neural Networks by Co-designing Activation Functions
[USENIX'24]	Abdulrahman Diaa, Lucas Fenaux, Thomas Humphries, Marian Dietz, Faezeh Ebrahimiaghazani, Bailey Kacsmar, Xinda Li, Nils Lukas , Rasoul Akhavan Mahdavi, Simon Oya, Ehsan Amjadian, Florian Kerschbaum. In the 33rd USENIX Security Symposium, 2024.
AR: 19.1% (417/2,176)	
9	Leveraging Optimization for Adaptive Attacks on Image Watermarks
[ICLR'24]	Nils Lukas , Abdulrahman Diaa, Lucas Fenaux, Florian Kerschbaum. In the Twelfth International Conference on Learning Representations, 2024.
AR: 30.8% (2,250/7,262)	
8	Universal Backdoor Attacks
[ICLR'24]	Benjamin Schneider, Nils Lukas , Florian Kerschbaum. In the Twelfth International Conference on Learning Representations, 2024.
AR: 30.8% (2,250/7,262)	
📰 Media Coverage	
7	PTW: Pivotal Tuning Watermarking for Pre-Trained Image Generators
[USENIX'23]	Nils Lukas and Florian Kerschbaum. In the 32nd USENIX Security Symposium, 2023.
AR: 29.2% (422/1,444)	
6	Analyzing Leakage of Personally Identifiable Information in Language Models
[S&P'23]	Nils Lukas , Ahmed Salem, Robert Sim, Shruti Tople, Lukas Wutschitz, Santiago Zanella-Béguelin. In the 44th IEEE Symposium on Security and Privacy, 2023.
AR: 17.0% (195/1,147)	
★ Distinguished Contribution Award at Microsoft MLADS	
5	SoK: How Robust is Image Classification Deep Neural Network Watermarking?
[S&P'22]	Nils Lukas , Edward Jiang, Xinda Li, Florian Kerschbaum. In the 43rd IEEE Symposium on Security and Privacy, 2022.
AR: 14.5% (147/1,012)	
4	Deep Neural Network Fingerprinting by Conferrable Adversarial Examples
[ICLR'21]	Nils Lukas , Yuxuan Zhang, Florian Kerschbaum. The Ninth International Conference on Learning Representations, 2021.
AR: 28.7% (860/2,997)	
★ Spotlight (Top 5%)	
3	On the Robustness of Backdoor-based Watermarking in Deep Neural Networks
[IH&MMSEC'21]	Masoumeh Shafieinejad, Nils Lukas , Jiaqi Wang, Xinda Li, Florian Kerschbaum. Proceedings of the 2021 ACM Workshop on Information Hiding and Multimedia Security, 2021.
AR: 40.3% (128/318)	
2	Practical Over-Threshold Multi-Party Private Set Intersection
[ACSAC'20]	Rasoul Mahdavi, Thomas Humphries, Bailey Kacsmar, Simeon Krastnikov, Nils Lukas , John Premkumar, Masoumeh Shafieinejad, Simon Oya, Florian Kerschbaum, Erik-Oliver Blass. Annual Computer Security Applications Conference (ACSAC), 2020.
AR: 20.9% (104/497)	
1	Differentially Private Two-Party Set Operations
[EuroS&P'20]	Bailey Kacsmar, Basit Khurram, Nils Lukas , Alexander Norton, Masoumeh Shafieinejad, Zhiwei Shang, Yaser Baseri, Maryam Sepehri, Simon Oya, Florian Kerschbaum. IEEE European Symposium on Security and Privacy (EuroS&P), 2020.
AR: 20.9% (39/187)	

PostDocs,
Students

Postdoctoral Researchers

• [Samuele Poppi](#)

since 2025

	• Kshitij Mishra (<i>secondary: Martin Takáč & Salem Lahlou</i>) since 2025 • Dushyant Chauhan (<i>secondary: Martin Takáč & Salem Lahlou</i>) since 2025 • Yuhan Liu (<i>secondary: Martin Takáč & Salem Lahlou</i>) since 2025
	Research Engineers
	• Sarim Hashmi since 2025
	PhD Students
	• Fabian Zhang (ML) since 2025 • Tair Djanibekov (NLP) since 2025 • Rushil Thareja (NLP) (<i>secondary: Praneeth Vepakomma</i>) since 2024 • Toluwani Aremu (ML) since 2023
	MSc Students
	• Juan Nicolas Sepulveda Arias (CV) since 2025 • Mohamed Mahmoud Mohamed Hendy (ML) since 2025 • Majid Mohamed Mohamed Sulaiman Ibrahim (ML) since 2025 • Prince Jha (ML) (<i>primary: Salem Lahlou</i>) since 2025 • Muhammad Mohideen (CV) (<i>secondary: Samuel Horvath</i>) since 2025
	Alumni
	• Tameem Bakr (ML) MSc · Mar 2026 • Shaikha Jasem Mohamed Ismail Alhosani (ML) MSc · Mar 2026 • Maryam Abdulla Rashed Abdulla Alshamsi (ML) MSc · Mar 2026 • Maryam Alshehyari (ML) MSc · Mar 2026 • Abdalla Khalid Mohamed Abdalla Alzaabi (ML) MSc · Mar 2026 • Ali Al-Ali (ML) MSc · Mar 2026
Talks & Keynotes	AI Safety Workshop: Alignment, Oversight and Deception
	• ADIA Lab & UGR Summer School, Granada 2026
	The Role of Watermarking for ML Security
	• CyberQ, UAE Cyber Security Council & TII, Abu Dhabi 2025
	Adaptively Robust and Forgery-Resistant Watermarking
	• Meta (FAIR), hosted by Hady Elsahar 2025 • International Symposium on Trustworthy Foundation Models, MBZUAI 2025
	Emerging Topics in Machine Learning
	• Build with AI, Google Developer Groups, Abu Dhabi 2025
	Optimizing Adaptive Attacks against Content Watermarks
	• DeepMind, hosted by David Stutz 2024 • University of California, Berkeley, hosted by Dawn Song 2024
	Analyzing Leakage of Personal Information in Language Models
	• Microsoft M365, hosted by Robert Sim 2024 • Meta, hosted by Will Bullock 2023 • MongoDB, hosted by Marilyn George and Archita Agarwal 2023
	How Reliable is Watermarking for Image Generators?
	• Google, hosted by Somesh Jha 2023 • University of California, Berkeley, hosted by Dawn Song 2023
	Keynotes
	• Aviation Future Week, hosted by Emirates, Dubai 2024 • Cyber Energy Leadership Forum, Abu Dhabi 2024
Honors & Awards	Amazon Research Award [100,000 USD] 2025
	First Place at the NeurIPS'24 Watermarking Competition [4,400 USD] 2024
	First Place at DGE Elite Hackathon, GITEX'24 [10,900 USD] 2024
	Top Mathematics Doctoral Thesis , University of Waterloo [1,080 USD] 2024

		Alumni Gold Medal, One PhD Award Yearly, University of Waterloo	2024
		Best Poster Award, Sponsored by David R. Cheriton [220 USD]	2023
		Distinguished Contribution Award, Microsoft MLADS conference	2023
		David R. Cheriton Scholarship, University of Waterloo [14,400 USD]	2022, 2023
		Outstanding Reviewer, ICML'22	2022
		Best Poster Award, Sponsored by Rogers [720 USD]	2019
Workshop Papers	3 [GenAI4Health]	Sanitizing Medical Documents with Differential Privacy using Large Language Models Rushil Thareja, Gautam Gupta, Preslav Nakov, Praneeth Vepakomma, Nils Lukas. 2025.	
	2 [WMARK'25]	First-Place Solution to NeurIPS 2024 Invisible Watermark Removal Challenge Fahad Shamshad, Tameem Bakr, Yahia Salaheldin Shaaban, Noor Hazim Hussein, Karthik Nandakumar and Nils Lukas. The 1st Workshop on GenAI Watermarking, 2025.	
	1 [WMARK'25] ★ Oral Presentation	Optimizing Adaptive Attacks against Content Watermarks for Language Models Abdulrahman Diaa, Toluwani Aremu, Nils Lukas. The 1st Workshop on GenAI Watermarking, 2025.	
Working Papers		Mitigating Watermark Forgery in Generative Models via Randomized Key Selection, Submitted. Toluwani Aremu, Noor Hazim Hussein, Munachiso S Nwadike, Samuele Poppi, Jie Zhang, Karthik Nandakumar, Neil Zhenqiang Gong, Nils Lukas. 2025.	
		Robust and Calibrated Detection of Authentic Multimedia Content, Submitted. Sarim Hashmi, Abdelrahman Elsayed, Mohammed Talha Alam, Samuele Poppi, Nils Lukas. 2025.	
Journal Publications	2 [IEEE S&P'23]	Privacy-Preserving Machine Learning [Cryptography] Florian Kerschbaum, Nils Lukas. IEEE Security & Privacy, Volume 21, Issue 6, 2023.	
	1 [AIP'18]	SunFlower: A new Solar Tower Simulation Method for use in Field Layout Optimization Pascal Richter, Gregor Heimig, Nils Lukas, Martin Frank. AIP Conference Proceedings, Volume 2033, Issue 1, 2018.	
Service	Track Chair		
		• ACM Conference on Computer and Communications Security (CCS), Security and Privacy of Machine Learning (co-chair)	2027
	Area Chair		
		• AAAI Conference on Artificial Intelligence (AAAI), Senior Program Committee	2027
		• International Conference on Learning Representations (ICLR)	2026, 2027
		• Neural Information Processing Systems (NeurIPS)	2026
		• International Conference on Machine Learning (ICML)	2026
	Program Committee		
		• ACM Conference on Computer and Communications Security (CCS)	2025
		• IEEE Symposium on Security and Privacy (IEEE S&P)	2025–2027
		• Recent Advances in Intrusion Detection (RAID)	2024
	Artifact Evaluation Committee		
		• The ACM Conference on Computer and Communications Security (CCS)	2023, 2024

Reviewer

- European Conference on Computer Vision (ECCV) 2026
- Conference on Language Modeling (COLM) 2026
- British Machine Vision Conference (BMVC) 2026
- FAGEN Workshop @ ICML 2026
- Workshop on GenAI Watermarking (WMark @ ICLR) 2025
- NETYS 2025
- ACM TheWebConf (WWW) 2025
- International Conference on Learning Representations (ICLR) 2024, 2025
- International World Wide Web Conference (TheWebConf) 2024
- Recent Advances in Intrusion Detection (RAID) 2023
- Neural Information Processing Systems (NeurIPS) 2022, 2023
- International Conference on Machine Learning (ICML) 2022, 2025
- The Conference on Information and Knowledge Management (CIKM) 2020

Journal Reviewer

- Communications of the ACM (CACM) 2025

Other

- External Examiner, PhD Dissertation, CISPA Helmholtz Center for Information Security 2026
- Faculty Search Committee, Machine Learning Department at MBZUAI 2025–2026
- Chair for the invited faculty talk program, International Symposium on Trustworthy Foundation Models at MBZUAI 2025
- Admission's Committee, MBZUAI Machine Learning Department 2025
- Student Board Member, Cybersecurity and Privacy Institute 2022–2024
- Session Chair, IEEE Symposium on Security and Privacy (S&P) 2023
- Sub-Reviewer, Proceedings on Privacy Enhancing Technologies (PETS) 2021–2023
- School Advisory Committee on Appointments Liaison, CrySP Lab 2022
- Organization, Workshop on Semantic Web Solutions for Large-Scale Biomedical Data Analytics (SeWeBMeDA) 2018

Teaching

Instructor, MBZUAI, UAE

- ML8502: Machine Learning Security (14 weeks) 2025, 2026
- ML8509: Collaborative Machine Learning (w/ S. Horvath) 2026
- ML807: Federated Learning (7 weeks) 2025
- ML818: Emerging Topics in Trustworthy Machine Learning (4 weeks) 2024

Teaching Assistant, University of Waterloo, Canada

- CS458/658: Computer Security and Privacy 2020, 2021
- CS246 – Object Oriented Programming 2021

Co-Instructor, RWTH-Aachen, Germany

- Course: Data-driven Medicine 2018

Education

University of Waterloo, Canada 2019 – 02/2024

Ph.D. in Computer Science

- Advisor: [Florian Kerschbaum](#)
- Thesis: Analyzing Threats of Large-Scale Machine Learning Systems
- Thesis Awards: [Top Mathematics Doctoral Prize](#) & [Alumni Gold Medal](#)

RWTH-Aachen, Germany

M.Sc. in Computer Science (*w/Distinction*) 2016 – 2018

B.Sc. in Computer Science 10/2012 – 2016

Work Experience

Assistant Professor, MBZUAI, Abu Dhabi, UAE since 08/2024

Research Intern, Royal Bank of Canada, Borealis AI, Toronto 2024

- Vertical Federated Learning, hosted by [Kevin Wilson](#)

Research Intern , Microsoft Research, Cambridge, UK	2022
• Privacy for Language Models, hosted by Shruti Tople & Lukas Wutschitz	
Research Assistant , RWTH-Aachen, Aachen	2014 – 2018
Student Researcher , DSA Daten- und Systemtechnik GmbH, Aachen	2016
Software Engineer Intern , A.R. Bayer DSP Systeme GmbH, Düsseldorf	2012